

The LegalBrain Indic Legal Corpus: A Foundational Resource for Generative AI in the Indian Judiciary

Anonymous CVPR Submission

Abstract

The integration of Artificial Intelligence into the legal domain holds transformative potential for the Indian judicial system, which currently grapples with a backlog of over 44 million cases. However, the development of robust “Legal AI” has been stymied by a critical scarcity of high-quality, large-scale training data specific to the Indian jurisprudential context. This report introduces and comprehensively analyzes the LegalBrain Indic Legal Corpus (hosted as Prarabdha/indian-legal-supervised-fine-tuning-data), a massive-scale dataset comprising approximately 6.06 million rows of Indian legal texts [8].

Unlike previous benchmarks that focused on narrow tasks like document segmentation or judgment prediction, this corpus is architected specifically for Supervised Fine-Tuning (SFT) of Large Language Models (LLMs), aiming to bridge the gap between general-purpose linguistic competence and specialized legal reasoning.

1. Introduction

1.1. The Jurisprudential Crisis and the AI Opportunity

The Indian legal system is one of the most extensive and complex in the world, serving a population of 1.4 billion people with a diverse array of linguistic, cultural, and socio-economic backgrounds. As of recent estimates, the National Judicial Data Grid indicates a pendency of over 44 million cases across the various tiers of the judiciary. This backlog results in significant delays, often summarized by the adage “justice delayed is justice denied.” The human capacity to process, analyze, and adjudicate this volume of litigation is stretched to its breaking point.

In this context, Artificial Intelligence (AI), particularly Generative AI driven by Large Language Models (LLMs), offers a compelling solution. The potential applications are vast: automated case summarization, intelligent legal research, drafting assistance for overburdened judges, and vernacular translation of complex judgments for litigants. However, the efficacy of such systems is entirely dependent

on the data upon which they are trained. A model trained primarily on Western legal datasets will fail to grasp the specific procedural intricacies of the Code of Civil Procedure, 1908, or the constitutional nuance of Article 21 as interpreted by the Supreme Court of India.

1.2. The “Cold Start” Problem in Indian Legal NLP

Until recently, the landscape of Indian Legal Natural Language Processing (NLP) suffered from a “cold start” problem. While general English datasets are abundant, domain-specific Indian legal data has been fragmented, inaccessible, or locked in proprietary databases. Researchers often relied on small, manually annotated datasets which, while high in quality, lacked the volume necessary to train the deep neural networks that power modern AI.

Existing datasets often focused on discriminative tasks—such as predicting whether a case would be allowed or dismissed—rather than generative tasks. For instance, the Indian Legal Documents Corpus (ILDC) provided 35,000 documents for judgment prediction [5]. , and OpenNyAI provided excellent resources for Named Entity Recognition (NER) and rhetorical role segmentation. However, to train an LLM to write like a lawyer or reason like a judge, one needs millions of tokens of continuous, diverse, and complex legal narrative.

1.3. The LegalBrain Indic Legal Corpus

The release of the LegalBrain Indic Legal Corpus (Prarabdha/indian-legal-supervised-fine-tuning-data) marks a paradigm shift in this ecosystem. With a volume of approximately 6.06 million rows, this dataset is orders of magnitude larger than its predecessors. It represents a move from “task-specific” data to “foundation-model-ready” data.

This report aims to:

1. Deconstruct the Dataset: Analyze the technical specifications, format (Parquet), and scale of the corpus.
2. Contextualize the Contribution: Rigorously compare the corpus against existing benchmarks like ILDC, OpenNyAI, and AI4Bharat to highlight the unfilled gaps it addresses.
3. Analyze Legal Substance: Perform a granular legal anal-

ysis of the text samples provided—covering administrative, family, and tax law—to illustrate the depth of reasoning the data contains.

4. **Propose Methodologies:** Outline how this data should be used for Supervised Fine-Tuning (SFT) to mitigate issues like hallucination and improve legal reasoning alignment.

2. The Pre-Existing Landscape: A Critical Review

To understand the monumental significance of a 6-million-row dataset, one must first appreciate the limitations of the datasets that came before. The Indian Legal NLP community has been active, but resource-constrained.

2.1. The OpenNyAI Mission: Structure over Scale

The OpenNyAI mission has been a pioneer in creating Digital Public Goods (DPGs) for the justice sector [7]. . Their philosophy has historically centered on structure—teaching machines to understand the anatomy of a legal document.

- **InLegalNER (Named Entity Recognition):** This dataset focuses on extracting specific entities such as “Petitioner,” “Respondent,” “Statute,” and “Precedent.” It contains approximately 46,545 annotated entities. While crucial for information extraction, it does not teach a model how to construct an argument, only how to identify the actors within one.
- **InRhetoricalRoles:** This corpus segments judgments into functional parts: Facts, Issues, Arguments, and Decision. It is highly curated but relatively small (<1,000 documents for some tasks). It is excellent for summarization pipelines but insufficient for training a base model to understand the content of the law.
- **Aalap:** OpenNyAI’s attempt at a generative model, Aalap was fine-tuned on a curated set of instructions derived from judgments and FIRs [6]. . While a significant step forward, the training data was focused on specific tasks like “statute retrieval” and “argument generation” based on a limited set of high-quality samples. It lacks the sheer exposure to the millions of variations in legal language that the LegalBrain corpus offers.

2.2. The Indian Legal Documents Corpus (ILDC): Predictive vs. Generative

The ILDC, introduced by Malik et al., was a landmark release containing roughly 35,000 Supreme Court cases.

- **Objective:** The primary goal of ILDC was Court Judgment Prediction (CJP)—predicting the outcome (acquittal/conviction) based on the facts.
- **Limitation:** While 35,000 is a respectable number for classification tasks, it is negligible for Language Modeling. An LLM trained on only 35,000 documents would

quickly overfit, memorizing the specific cases rather than learning the general principles of law. Furthermore, ILDC is restricted to Supreme Court cases, missing the vast and chaotic reality of High Courts and Tribunals where the bulk of litigation occurs.

2.3. AI4Bharat: The Multilingual Frontier

AI4Bharat has done exceptional work in democratizing access to Indic languages [1, 2]. through datasets like BPCC (Bharat Parallel Corpus Collection) and Samanantar.

- **Focus:** Their primary focus is Machine Translation (MT) and Automatic Speech Recognition (ASR). The BPCC contains approximately 230 million parallel text pairs.
- **Relevance:** While they have released legal glossaries, their datasets are designed to teach a computer that “lawyer” in English equals “vakeel” in Hindi. They are not primarily designed to teach the computer what a lawyer does or the legal arguments regarding “anticipatory bail” under Section 438 of the CrPC. The LegalBrain corpus fills this specific niche of substantive legal knowledge rather than just linguistic translation.

2.4. Comparative Summary

The following table summarizes the positioning of the LegalBrain Indic Legal Corpus against established benchmarks.

3. Dataset Architecture and Technical Specifications

The utility of a dataset is defined not just by its content, but by its accessibility and technical format. The LegalBrain corpus utilizes modern data engineering standards to ensure it is “AI-ready.”

3.1. Repository and Access

The dataset is hosted on Hugging Face under the repository ID Prarabdha/indian-legal-supervised-fine-tuning-data.

- **Visibility:** It is a public dataset, accessible to the global research community.
- **License:** It carries an Apache 2.0 license. Unlike Creative Commons Non-Commercial licenses, Apache 2.0 permits commercial use. This allows legal-tech startups to build proprietary tools using this open data, fostering an ecosystem of innovation rather than just academic research.

3.2. File Format and Efficiency

The data is stored in Parquet format.

- **Columnar Storage:** Unlike JSON or CSV, which are row-based, Parquet is column-based. This allows for highly efficient compression and faster retrieval of specific data fields. For a dataset approaching 10 GB, this

Table 1. Comparative Analysis of Major Indian Legal NLP Datasets

Feature	LegalBrain Corpus	Indic Legal	OpenNyAI	ILDC	AI4Bharat
Primary Task	Generative AI / SFT / Reasoning		Entity Extraction / Segmentation	Judgment Prediction	Translation / ASR
Size	~6,060,000 rows		~10k–46k entities	~35,000 docs	~230M pairs
Content Depth	Full text, reasoning, citations		Annotated spans, segments	Facts mapped to outcomes	Parallel sentences, terms
Target Model	LLMs (LLaMA, Mistral)		BERT, RoBERTa	Hierarchical Attention Nets	NLLB, IndicTrans
Legal Scope	Supreme Court, High Courts, Tribunals		Supreme Court focus	Supreme Court focus	General + Legal Terms

efficiency is vital for data loading pipelines during the training of large models.

- **Integration:** The dataset card indicates support for libraries like Dask and Croissant.
 - Dask enables parallel computing, allowing researchers to process this 6-million-row dataset on distributed clusters rather than loading it all into a single machine’s RAM.
 - Croissant is a metadata format that makes datasets more discoverable and consumable by ML tools, indicating adherence to modern MLOps standards.

3.3. Data Volume and Splits

- **Total Rows:** Approximately 6.06 million.
- **Splits:** The dataset includes a train split of 6.06M rows. Researchers are expected to carve out their own validation and test sets, or use the separate indic-legal-benchmark dataset for evaluation.
- **Size:** The disk footprint is approximately 9.1 GB. This density suggests that the text per row is substantial, likely containing full paragraphs or pages of judgments.

4. Substantive Legal Analysis: Decoding the Data

The true value of the LegalBrain corpus lies in the content of the text. It is not merely a collection of words; it is a repository of judicial reasoning. By analyzing specific snippets provided in the dataset documentation, we can reconstruct the types of legal logic an AI model would learn from this data.

4.1. Administrative Law: The Principle of Equity

Source: Snippets

One of the detailed examples in the dataset involves a dispute between the Delhi Development Authority (DDA) and land allottees (Case: D.D.A v. K.K. Mehta).

The Legal Fact Pattern:

The DDA allotted land to applicants but could not hand over possession immediately due to a stay order from the High Court regarding the land acquisition. Years later, when the stay was lifted, the DDA demanded “escalation charges” (higher prices) from the allottees. The allottees challenged this, arguing they should not pay for delays they did not cause.

The Judicial Reasoning (The “Instruction” for the AI):

The text captures a sophisticated argument based on the Principle of Equity.

1. **Issue:** Can DDA charge escalation for a delay caused by judicial intervention?
2. **Ratio:** The High Court held that if DDA charges escalation, it must “equitably” adjust the interest. Specifically, it should pay interest on the deposit money the allottees had locked up with the DDA for those years.
3. **Outcome:** The Court directed the DDA to charge escalation at the same rate of interest that it pays to unsuccessful applicants, creating a financial balance.

AI Learning Outcome:

An LLM trained on this learns:

- Concept: “Equitable Set-off.”
- Context: Administrative delays and “Stay Orders.”
- Causality: If a delay is caused by a Court Order (Vis Major / Act of State equivalent), the private party should not be penalized.
- Vocabulary: “Escalation charges,” “Allotment letter,” “Draw of plots.”

4.2. Family Law and Bio-Ethics: The DNA Paternity Conflict

Source: Snippets

A highly sensitive and complex area of law covered in the corpus is the use of DNA tests to determine paternity in

matrimonial disputes. This section is critical for teaching AI how to handle conflicting precedents.

The Conflict:

The dataset explicitly contrasts two landmark Supreme Court judgments:

1. **Gautam Kundu vs. State of West Bengal:** Established a conservative precedent. Courts cannot order blood tests “as a matter of course.” There is a strong presumption of legitimacy of a child born during marriage (Section 112, Evidence Act). The court must avoid “roving inquiries.”
2. **Sharda vs. Dharmpal:** A later judgment that softened this stance. It held that a Matrimonial Court has the power to order medical tests if there is a “strong prima facie case” and sufficient material.

The Synthesis:

The dataset text provides the synthesis: “Is it permissible for a court to order a DNA test... Yes... if it is eminently needed, after balancing the interests of the parties involved.”

AI Learning Outcome:

- **Nuance:** The model learns that the answer to “Can I get a DNA test?” is not “Yes” or “No,” but “It depends on eminent need.”
- **Citation Usage:** It learns to cite Gautam Kundu when arguing for privacy and presumption of legitimacy, and Sharda when arguing for the truth and medical evidence.
- **Statutory Knowledge:** It links these arguments to Section 112 of the Indian Evidence Act.

4.3. Corporate Taxation: Statutory Interpretation

Source: Snippets

The dataset includes technical interpretation of the Income Tax Act, 1961.

The Provision:

Section 80I allows deductions for profits from industrial undertakings.

The Condition:

The text clarifies that for a unit to claim this, it must be an “independent industrial undertaking.”

The Evidence Requirement:

The judgment cites *CIT vs. Mount Shivalik Breweries Ltd.* to establish that “independence” is proven by maintaining separate books of account, specifically a separate Balance Sheet and Profit & Loss Account.

AI Learning Outcome:

- **Compliance Logic:** The model learns a functional rule: “Want Section 80I deduction? → Must keep separate books.”
- **Specifics:** It learns that “books” specifically means Balance Sheet and P&L Account, not just a ledger.

5. Methodology: Supervised Fine-Tuning (SFT) for Legal AI

The release of 6 million rows of such data enables a specific training methodology known as Supervised Fine-Tuning (SFT). This is distinct from “Pre-training.”

5.1. From Pre-training to SFT

- **Pre-training:** Teaches a model grammar and general world knowledge (for example, “The sky is blue”). Most models are pre-trained on the general internet. They know about law, but they do not know Indian law deeply.
- **SFT:** Uses a curated dataset such as LegalBrain to teach the model how to respond to specific prompts or follow specific instructions in a specific domain.

5.2. Addressing Hallucination

A significant risk in Legal AI is hallucination—the invention of fake cases. The dataset documentation ironically mentions a tool called “LegalBrain” that “provides a well-articulated paragraph, citing several seemingly authoritative but, in reality, fabricated cases.” This highlights the danger of using general-purpose models for law [3].

The Solution:

By fine-tuning on the LegalBrain Corpus, we saturate the model’s parameters with real case names such as DDA v. K.K. Mehta and Sharda v. Dharmpal.

- **Probability Shift:** When the model tries to predict the next word after “Supreme Court judgment on DNA tests,” the probability of it predicting “Sharda vs. Dharmpal” becomes significantly higher than predicting a hallucinated name, because it has seen “Sharda vs. Dharmpal” thousands of times in the training data.
- **Grounding:** The dataset serves as the ground truth. SFT aligns the model’s output distribution with the actual distribution of Indian legal texts.

5.3. Instruction Tuning Possibilities

The “supervised-fine-tuning” label suggests the data is likely structured or can be structured into instruction–output pairs.

- **Instruction:** “Summarize the court’s view on escalation charges.”
- **Output:** Generated summary grounded in precedent.

Researchers can use this to build agentic workflows. For example, the Aalap model showed that training on specific tasks such as “Issue Generation” improves performance. The LegalBrain corpus allows this to be done at one hundred times the scale.

6. Comparison with Emerging Datasets (2024–2025)

The legal data space is dynamic. Recent months have seen other releases, yet LegalBrain remains dominant in scale.

- **IndicLegalQA (2024):** A dataset of 10,000 question–answer pairs from Supreme Court judgments [11].
 - **Comparison:** 10,000 is useful for evaluation, but too small for training. LegalBrain’s 6 million rows provide the “textbook,” while IndicLegalQA provides the “quiz.”
- **IndianBailJudgments (2024):** A specialized dataset of 1,200 bail orders annotated with GPT-4 [10].
 - **Comparison:** Highly specialized for one domain (bail). LegalBrain covers all domains including civil, tax, family, and constitutional law.
- **SC Judgments (Kaggle):** A raw dump of 26,000 PDF files [4].
 - **Comparison:** Raw PDFs require OCR and cleaning. LegalBrain provides Parquet-formatted text that is clean and ready for ingestion, saving months of pre-processing time.

7. Societal Impact: Democratizing Legal Intelligence

The implications of this dataset extend beyond computer science into the realm of social justice.

7.1. Access to Justice

In India, high-quality legal representation is expensive. A model fine-tuned on LegalBrain could power a “Legal Assistant” for legal aid clinics in rural India.

Scenario: A paralegal in a village enters a query: “The government took my client’s land but hasn’t paid for 5 years.”

AI Response: Drawing on the DDA cases in the corpus, the AI can suggest arguments regarding “delay in possession” and “interest on deposits,” enabling the paralegal to draft a competent representation.

7.2. The Digital Divide and Language

While the current corpus is primarily English (the language of the Higher Judiciary), it sets the stage for multilingual models. By aligning this English legal reasoning with the translation capabilities of AI4Bharat’s models, future systems can reason in English while communicating in Hindi, Tamil, or Telugu.

8. Technical Roadmap and Future Directions

The release of the LegalBrain Indic Legal Corpus is a milestone, but not the finish line. The dataset documentation hints at future technical integrations.

8.1. Quantization and Edge Deployment

Community discussions around the dataset author mention using quantization and compression to run models on consumer-grade GPUs.

A 6-million-row dataset can train a 70 billion parameter model, but running such a model requires expensive servers. Techniques like 4-bit quantization (QLoRA) allow these powerful models to run on standard laptops, decentralizing legal AI.

8.2. Retrieval Augmented Generation (RAG)

While SFT teaches the model reasoning, specific facts change. The dataset documentation mentions using vector databases and caching systems.

The future lies in a hybrid approach: use the LegalBrain Corpus to SFT the model (teaching it how to think like a lawyer), then use RAG to fetch the latest judgments to ensure factual currency.

8.3. Benchmarking

The existence of the Prarabdha/indic-legal-benchmark dataset is crucial. It is not enough to claim improved performance; models must be evaluated against unseen benchmark data [9].

9. Conclusion

The LegalBrain Indic Legal Corpus is a foundational asset for the Indian Legal AI ecosystem. By providing 6.06 million rows of diverse, complex, and machine-readable legal text, it addresses the critical data scarcity that has long hindered the development of domain-specific LLMs for Indian law.

The dataset captures intricate reasoning required for administrative equity disputes, complex family law conflicts involving bio-ethics, and statutory compliance in corporate taxation. It moves beyond the limitations of previous datasets to enable generative legal reasoning.

As researchers and developers fine-tune models on this corpus, we anticipate a reduction in AI hallucinations and a significant improvement in automated legal research and drafting. In a country with over 44 million pending cases, such efficiency gains are not merely a technological convenience but a necessity.

References

- [1] AI4Bharat. Samanantar: The largest publicly available parallel corpus for indic languages, 2021.
- [2] AI4Bharat. Bharat parallel corpus collection (bpcc), 2023.
- [3] Ilias Chalkidis et al. Natural language processing for the legal domain: A survey of tasks, datasets, models and challenges. *arXiv*, 2024.
- [4] Kaggle. Sc judgments india dataset, 2023.

- [5] TIB Leibniz. Indian legal documents corpus dataset, 2023.
- [6] OpenNyAI. Aalap instruction dataset, 2024.
- [7] OpenNyAI. Opennyai project, 2024.
- [8] Prarabdha. Indian legal supervised fine-tuning dataset, 2026.
- [9] Prarabdha. Prarabdha hugging face profile, 2026.
- [10] ResearchGate. Indianbailjudgments-1200 dataset, 2024.
- [11] A. Singh et al. Dataset for legal question answering in the indian judiciary context. *PMC*, 2024.